

Neural Networks (part 1)

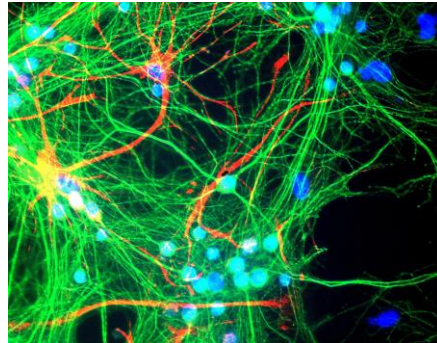
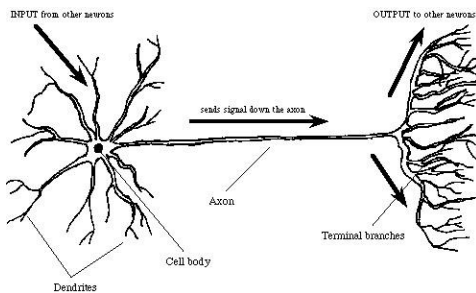
Goals for this lecture

you should understand the following concepts

- perceptrons
- the perceptron training rule
- linear separability
- hidden units
- multilayer neural networks
- gradient descent
- stochastic (online) gradient descent
- activation functions
- sigmoid, hyperbolic tangent, ReLU
- objective (error, loss) functions
- squared error, cross entropy

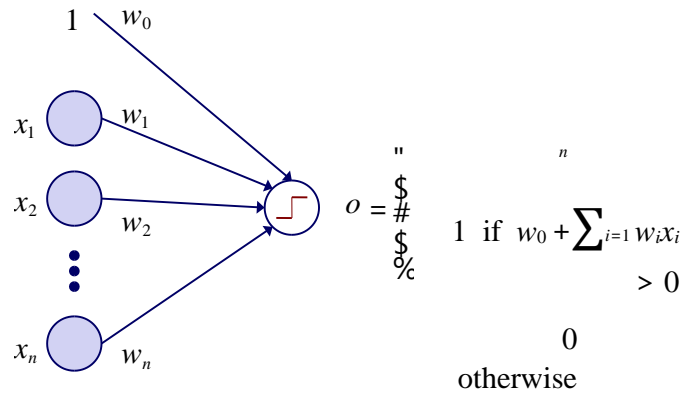
Neural Networks

- a.k.a. *artificial neural networks, connectionist models*
- inspired by interconnected neurons in biological systems
- simple processing units
- each unit receives a number of real-valued inputs
- each unit produces a single real-valued output



Perceptrons

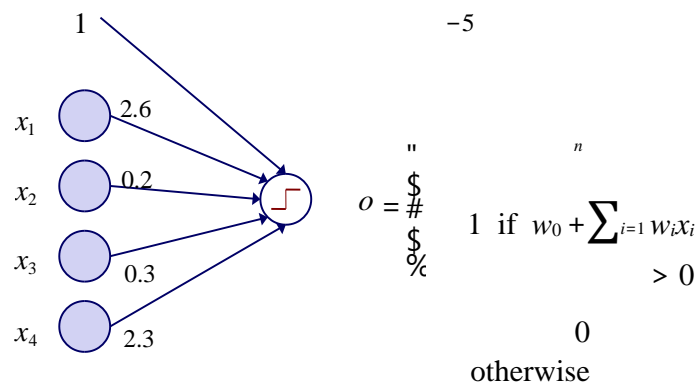
[McCulloch & Pitts, 1943; Rosenblatt, 1959; Widrow & Hoff, 1960]



input units:
represent given x

output unit:
represents binary classification

Perceptron example



features, class labels are represented numerically

$$\begin{aligned} \langle \quad \rangle & \quad \mathbf{x} = 1, 0, 0, 1 \quad w_0 + \sum_{i=1}^n w_i x_i = -0.1 \quad o = 0 \\ \langle \quad \rangle & \quad \end{aligned}$$

$$\mathbf{x} = 1, 0, 1, 1 \quad w_0 + \sum_{i=1}^n w_i x_i = 0.2 \quad o = 1$$

Training a perceptron

1. randomly initialize weights
2. iterate through training instances until convergence

2a. calculate the output for the given instance

$$o = \begin{cases} 1 & \text{if } w_0 + \sum_{i=1}^n w_i x_i > 0 \\ 0 & \text{otherwise} \end{cases}$$

2b. update each weight

$$\Delta w_i = \eta (y - o) x_i$$

η is *learning rate*;

set to value $\ll 1$

$$w \leftarrow w_i + \Delta w_i$$

Representational power of perceptrons

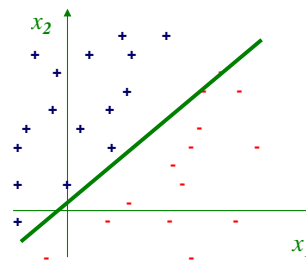
perceptrons can represent only *linearly separable* concepts

$$o = \begin{cases} 1 & \text{if } w_0 + \sum_{i=1}^n w_i x_i > 0 \\ 0 & \text{otherwise} \end{cases}$$

decision boundary given by: 1 if w_0

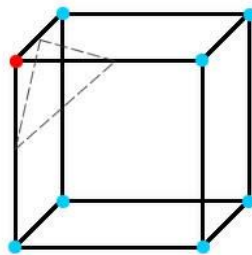
$$+w_1x_1 + w_2x_2 > 0 \quad w_1x_1 + w_2x_2 = -w_0$$

$$x_2 = -\frac{w_1x_1 - w_0}{w_2}$$



Representational power of perceptrons

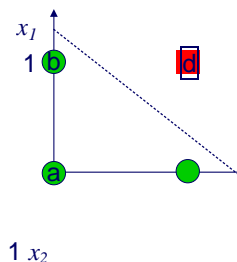
- in previous example, feature space was 2D so decision boundary was a line
- in higher dimensions, decision boundary is a hyperplane



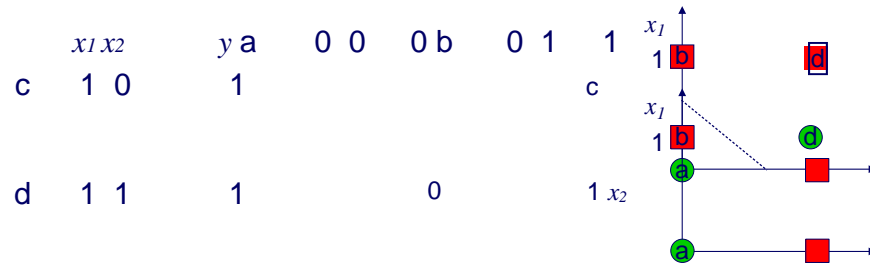
Linearly separable functions

AND

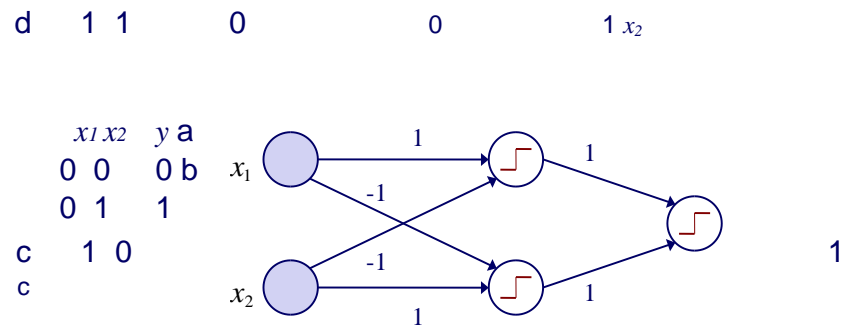
	x_1, x_2	y	a	b	c	d
0						
b	0 1	0				
c	1 0	0		c		
d	1 1	1	0			



OR



Is XOR linearly separable?



a multilayer perceptron
can represent XOR

assume $w_0 = 0$ for all nodes

Example of multi-layer neural network

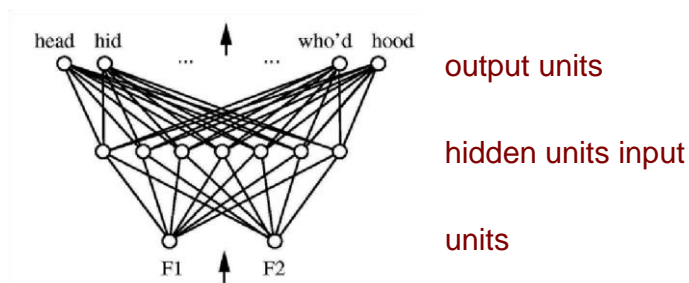


figure from Huang & Lippmann, *NIPS* 1988

input: two features from spectral analysis of a spoken sound output:
vowel sound occurring in the context “h__d”

Decision regions of a multi-layer NN

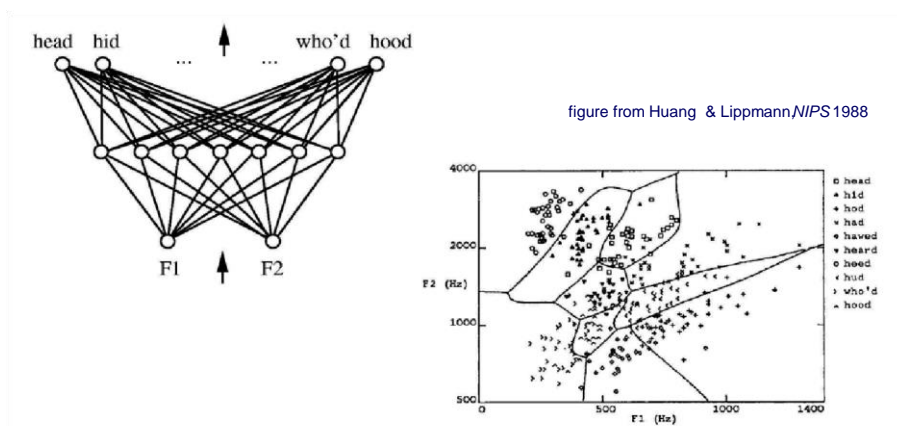
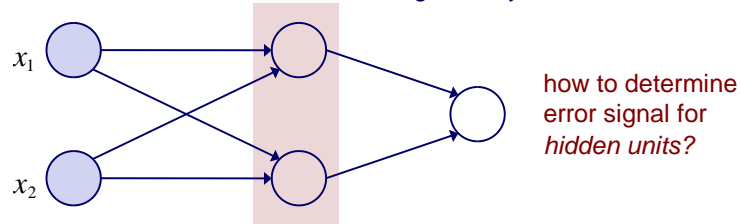


figure from Huang & Lippmann, *NIPS* 1988

input: two features from spectral analysis of a spoken sound output:
vowel sound occurring in the context “h__d”

Learning in a multi-layer NN

- work on neural nets fizzled in the 1960's
- single layer networks had representational limitations (linear separability)
- no effective methods for training multilayer networks



- revived again with the invention of *backpropagation* method [Rumelhart & McClelland, 1986; also Werbos, 1975]
- key insight: require neural network to be differentiable; use *gradient descent*

Gradient descent in weight space

Given a training set

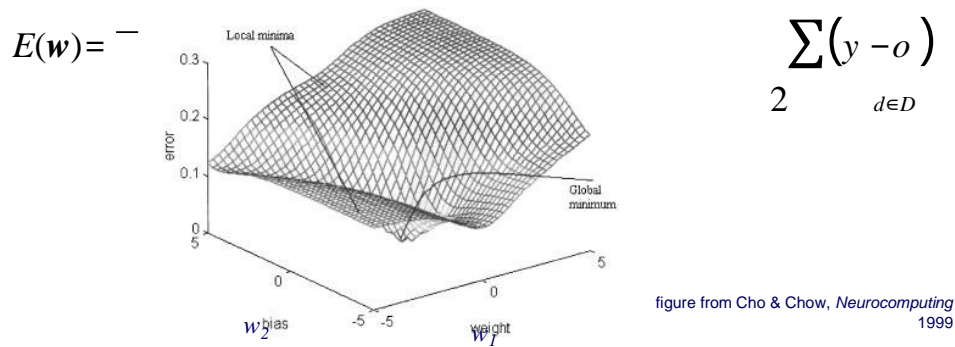
we can specify an

$$D = \{(\mathbf{x}^{(1)}, y^{(1)}) \square (\mathbf{x}^{(m)}, y^{(m)})\}$$

error measure that is a function of our weight

vector \mathbf{w}

$$1 \quad (d) \quad (d)^2$$



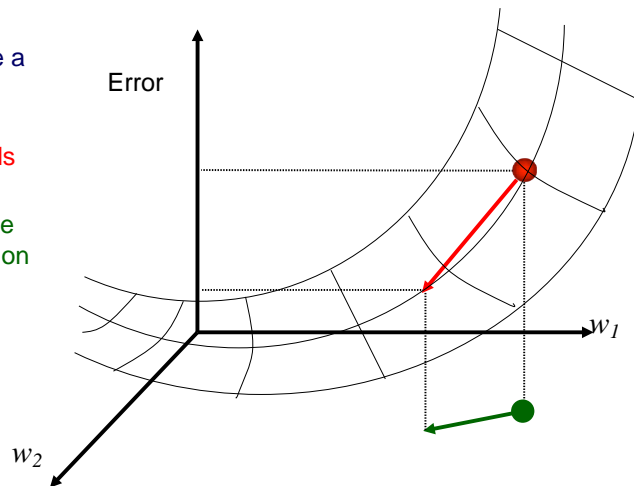
This *objective function* defines a surface over the model (i.e. weight) space

Gradient descent in weight space

gradient descent is an iterative process aimed at finding a minimum in the error surface

on each iteration

- current weights define a point in this space
- find direction in which error surface descends most steeply
- take a step (i.e. update weights) in that direction



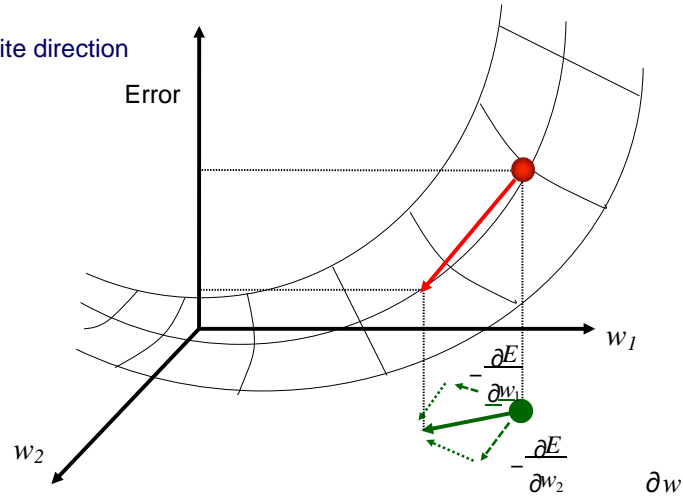
Gradient descent in weight space

calculate the gradient of E : $\nabla E(\mathbf{w}) = \left(\frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2}, \dots, \frac{\partial E}{\partial w_n} \right)$

take a step in the opposite direction

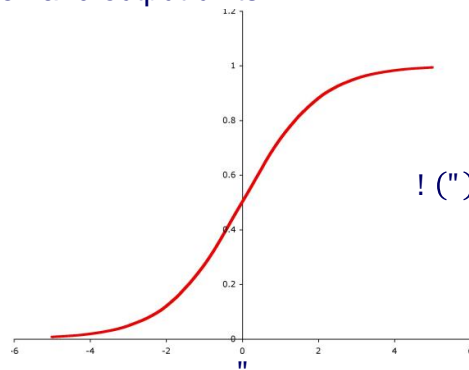
$$\Delta \mathbf{w} = -\eta \nabla E(\mathbf{w})$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$



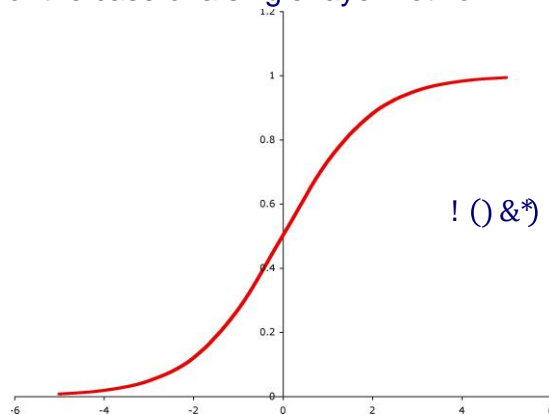
The sigmoid function

- to be able to differentiate E with respect to w_i , our network must represent a continuous function
- to do this, we can use *sigmoid functions* instead of threshold functions in our hidden and output units



The sigmoid function

for the case of a single-layer network



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z = 2 \cdot 3 + 4 \cdot \frac{2}{5} = 2.8$$

BatchNN training

given: network structure and a training set $D = \{(\mathbf{x}^{(1)}, y^{(1)}) \dots (\mathbf{x}^{(m)}, y^{(m)})\}$

initialize all weights in \mathbf{w} to small random numbers

until stopping criteria met do initialize the error

$$E(\mathbf{w}) = 0$$

for each $(\mathbf{x}^{(d)}, y^{(d)})$ in the training set input $\mathbf{x}^{(d)}$ to the network and compute output $o^{(d)}$ increment the error

$$E(\mathbf{w}) = E(\mathbf{w}) + \frac{1}{2} (y^{(d)} - o^{(d)})^2 \quad \text{calculate the gradient}$$

$$\nabla E(\mathbf{w}) = \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

update the weights

$$\Delta \mathbf{w} = -\eta \nabla E(\mathbf{w})$$

Online vs. Batch learning

- Standard gradient descent (batch training): calculates error gradient for the entire training set, before taking a step in weight space
- *Stochastic gradient descent* (online training): calculates error gradient for a single instance (or a small set of instances, a “mini batch”), then takes a step in weight space
 - much faster convergence
 - less susceptible to local minima

OnlineNN training (Stochastic Gradient Descent)

given: network structure and a training set $D = \{(\mathbf{x}^{(1)}, y^{(1)}) \sqcup (\mathbf{x}^{(m)}, y^{(m)})\}$

initialize all weights in \mathbf{w} to small random numbers until
stopping criteria met do

for each $(\mathbf{x}^{(d)}, y^{(d)})$ in the training set input $\mathbf{x}^{(d)}$ to the
network and compute output $o^{(d)}$ calculate the

error $E(\mathbf{w}) = \frac{1}{2} (y^{(d)} - o^{(d)})^2$ calculate the gradient

$$\nabla E(\mathbf{w}) = \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

update the weights

$$\Delta \mathbf{w} = -\eta \nabla E(\mathbf{w})$$

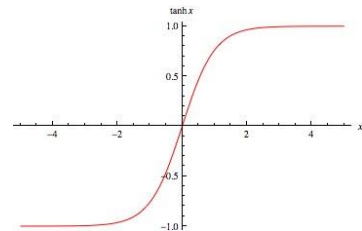
Other activation functions

- the sigmoid is just one choice for an *activation function*

there are others we can use including

hyperbolic tangent

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



rectified linear (ReLU)

$$\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

Other objective functions

- squared error is just one choice for an *objective function*

there are others we can use including cross entropy

$$J_{\text{CE}} = - \sum_{H \in \mathcal{N}} \sum_{G \in \mathcal{H}} y_i^{(d)} \ln(o_i^{(d)}) - (1 - y_i^{(d)}) \ln(1 - o_i^{(d)})$$


multiclass

cross entropy

$$J_{\text{CE}} = - \sum_{H \in \mathcal{N}} \sum_{G \in \mathcal{H}} y_i^{(d)} \ln(o_i^{(d)}) - (1 - y_i^{(d)}) \ln(1 - o_i^{(d)})$$

Convergence of gradient descent

- gradient descent will converge to a minimum in the error function
- for a multi-layer network, this may be a *local minimum* (i.e. there may be a “better” solution elsewhere in weight space)



- for a single-layer network, this will be a global minimum (i.e. gradient descent will find the “best” solution)
- Recent analysis suggests that local minima are probably rare in high dimensions; saddle points are more of a challenge [Dauphin et al., NIPS 2014]